

**Srinivas Prasad Prabhu**

**UT EID - sp55629**

**Machine Used – Intel Xeon – 16 cores, 2 GHz + 64 GB RAM**

**CS380P – Lab 1 – Prefix Scan and Barriers**



1) Recall that prefix scan requires a barrier for synchronization. In this step, you will write a work-efficient parallel prefix scan using pthread barriers. For each of the provided input sets, set the number of loops of the operator to 100000 (-l 100000) and graph the execution time of your parallel implementation over a sequential prefix scan implementation as a function of the number of worker threads used. Vary from 2 to 32 threads in increments of 2. Then, explain the trends in the graph. Why do these occur? From here: “A prefix sum algorithm is work-efficient if it does asymptotically no more work (add operations, in this case) than the sequential version. In other words, the two implementations should have the same work complexity,  $O(n)$ .”

Ans –

Please find below the graph of the execution times of various inputs with -l 10E5 vs varying number of worker threads from 0-32. The 2 graphs below plot the same values of microseconds in linear and logarithmic scale.

Parallel Prefix Scan execution time (usec) vs Number of Threads - Loops - 100000

| usec                   | 0      | 1      | 2      | 4      | 6      | 8      | 10     | 12     | 14     | 16     | 18     | 20     | 22     | 24     | 26     | 28     | 30     | 32      |
|------------------------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|---------|
| seq_64_test.txt/100000 | 1846   | 3905   | 3181   | 2937   | 3048   | 3134   | 3301   | 3705   | 3845   | 3846   | 4503   | 4312   | 5074   | 5187   | 5026   | 5343   | 5224   | 5364    |
| 1k.txt/100000          | 29877  | 60600  | 47432  | 41091  | 40924  | 41121  | 42503  | 43545  | 48267  | 49412  | 52285  | 55907  | 54669  | 55228  | 64604  | 66061  | 62941  | 69824   |
| 8k.txt/100000          | 239175 | 485217 | 374370 | 326198 | 320291 | 324428 | 333373 | 343062 | 352918 | 360557 | 380587 | 387800 | 397998 | 415162 | 438990 | 466699 | 482699 | 506798  |
| 16k.txt/100000         | 478240 | 969306 | 750513 | 650009 | 637764 | 644541 | 663123 | 682422 | 699517 | 722558 | 745963 | 773666 | 802446 | 823747 | 862904 | 915609 | 965021 | 1002646 |



## Observations -

- 1) It can be observed from the graph that a parallel implementation with just 1 thread is highly in-efficient and consumes exponentially higher time. This is because of the overheads involved in thread creation and the context switch time at each of the barriers.
- 2) As the number of threads increases, the work starts getting more evenly divided among the threads and we can see exponential reduction in times compared to a single thread.
- 3) From the logarithmic graph it can be observed that beyond 8-10 threads, the time consumed starts increasing higher than the sequential implementation even though the growth is slow.
- 4) As the number of threads increases, the size of the blocks of input on which each thread works on decreases. This leads to very high rate of cache invalidations. This reduces the actual speedup obtained by the parallel implementation.
- 5) As the number of threads increases beyond the number of physical cores, a lot of time is spent in context switching between the various threads. This again leads to reduction in the speedup obtained by the parallel implementation.
- 6) As all the worker threads are working on the same data, the need for barrier synchronization also reduces the performance of the parallel implementation as multiple threads are waiting for longer at the same barrier.
- 7) Due to points 4-6 above just increasing the number of threads doesn't yield a high-performance solution. Partitioning the data optimally so that different threads are working on different independent blocks of data (which fit into a cache line), which in turn reduces the need for barrier synchronization will improve the performance of parallel prefix scan.

2) Now that you have a working and, hopefully, efficient parallel implementation of prefix scan, try changing the amount of loops the operator does to 10(-l 10) and plot the same graph as before. What happened and why? Vary the -l argument and find the inflexion point, where sequential and parallel elapsed times meet (does not have to be exact, just somewhat similar). Why can changing this number make the sequential version faster than the parallel? What is the most important characteristic of it that makes this happen? Argue this in a general way, as if different -l parameters were different operators, afterall, the -l parameter is just a way to quickly create an operator that does something different.



Ans –

Please find below the graph of the execution times of various inputs with -l 10 vs varying number of worker threads from 0-32. The 2 graphs below plot the same values of microseconds in linear and logarithmic scale.

Second, graphs - various inputs with different loop numbers from 1 – 10E7 vs execution time is plotted on a logarithmic scale of time, with different fixed number of threads.



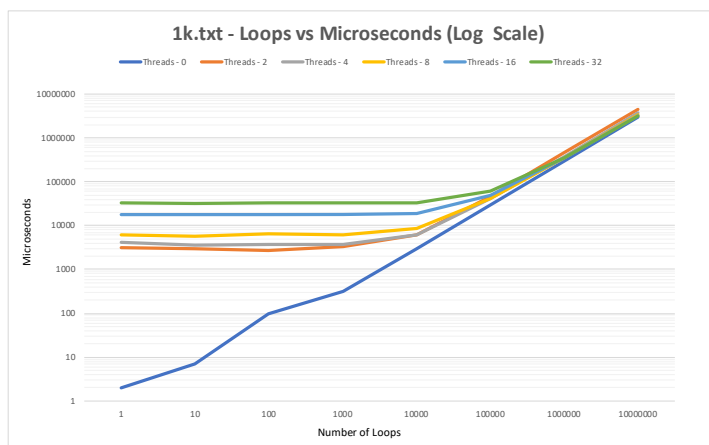

Parallel Prefix Scan execution time (usec) vs Number of Threads - Loops - 10

| usec               | 0   | 1    | 2     | 4     | 6     | 8     | 10     | 12     | 14     | 16     | 18     | 20     | 22     | 24     | 26     | 28     | 30     | 32     |
|--------------------|-----|------|-------|-------|-------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| seq_64_test.txt/10 | 0   | 170  | 532   | 705   | 961   | 1091  | 1178   | 1945   | 1540   | 1692   | 2437   | 2131   | 2242   | 2465   | 2675   | 2821   | 3000   | 3203   |
| 1k.txt/10          | 7   | 331  | 2919  | 3969  | 4780  | 5609  | 13233  | 14084  | 16373  | 17768  | 20154  | 21696  | 23194  | 25540  | 26997  | 28518  | 29932  | 32157  |
| 8k.txt/10          | 55  | 1366 | 20580 | 26758 | 35285 | 41365 | 102473 | 110888 | 125026 | 136001 | 153684 | 168028 | 177813 | 189261 | 204913 | 214366 | 228584 | 244994 |
| 16k.txt/10         | 111 | 2419 | 38917 | 53714 | 68906 | 85955 | 203071 | 224326 | 251215 | 279217 | 305523 | 330061 | 354790 | 386032 | 401802 | 428142 | 462205 | 487047 |



Execution time (usec) of different inputs with varying loops and fixed number of threads - pthread barrier

| seq_64_test.txt |             |             |             |             |              |              | 1k.txt   |             |             |             |             |              |              |
|-----------------|-------------|-------------|-------------|-------------|--------------|--------------|----------|-------------|-------------|-------------|-------------|--------------|--------------|
|                 | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |          | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |
| 1               | 0           | 676         | 749         | 1184        | 1970         | 3639         | 1        | 2           | 3121        | 4057        | 6118        | 18243        | 32389        |
| 10              | 0           | 554         | 786         | 955         | 2048         | 3530         | 10       | 7           | 2925        | 3595        | 5727        | 18182        | 32065        |
| 100             | 3           | 561         | 691         | 1196        | 2047         | 3391         | 100      | 97          | 2713        | 3625        | 6449        | 18053        | 32473        |
| 1000            | 20          | 620         | 812         | 1278        | 2069         | 2891         | 1000     | 312         | 3404        | 3800        | 6164        | 17977        | 32897        |
| 10000           | 186         | 676         | 900         | 1351        | 2256         | 3607         | 10000    | 3020        | 6147        | 6266        | 8589        | 19274        | 33044        |
| 100000          | 1838        | 3236        | 2910        | 3188        | 3608         | 5541         | 100000   | 29933       | 47451       | 41317       | 41232       | 48316        | 62250        |
| 1000000         | 18427       | 28080       | 23944       | 22155       | 21895        | 28541        | 1000000  | 298449      | 450203      | 377349      | 342555      | 338415       | 346548       |
| 10000000        | 183823      | 277998      | 234599      | 214082      | 208098       | 210820       | 10000000 | 2984304     | 4483120     | 3737520     | 3369506     | 3189994      | 3157642      |
| 8k.txt          |             |             |             |             |              |              | 16k.txt  |             |             |             |             |              |              |
|                 | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |          | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |
| 1               | 19          | 20388       | 27213       | 42133       | 141403       | 245891       | 1        | 39          | 42572       | 56311       | 85131       | 273836       | 485115       |
| 10              | 56          | 20314       | 27255       | 42300       | 140593       | 242263       | 10       | 111         | 40822       | 55396       | 83718       | 270694       | 486128       |
| 100             | 862         | 21793       | 27593       | 42775       | 141141       | 241676       | 100      | 1694        | 41313       | 54693       | 85090       | 271880       | 481648       |
| 1000            | 2540        | 23911       | 28847       | 43586       | 141103       | 243439       | 1000     | 5084        | 49517       | 56945       | 89148       | 276208       | 480844       |
| 10000           | 24065       | 47271       | 46686       | 59236       | 153481       | 258310       | 10000    | 48159       | 90224       | 94689       | 114670      | 302259       | 596191       |
| 100000          | 239121      | 376196      | 328149      | 324021      | 363584       | 510425       | 100000   | 478251      | 746777      | 654386      | 644787      | 723187       | 1013534      |
| 1000000         | 2389581     | 3599748     | 3012040     | 2734501     | 2618149      | 2698534      | 1000000  | 4779271     | 7193906     | 6018217     | 5508610     | 5240268      | 5399959      |
| 10000000        | 23893812    | 35850970    | 29907518    | 26967498    | 25475604     | 25129539     | 10000000 | 47790761    | 71721033    | 59777722    | 54195151    | 50935196     | 50020808     |



## Observations –

- 1) Just like previously seen, it can be observed from the graph that a parallel implementation with just 1 thread is highly in-efficient and consumes exponentially higher time. This is because of the overheads involved in thread creation and the context switch time at each of the barriers.
- 2) Sequential implementation seems to be much faster than parallel implementation for low values of loops. This is because the parallel implementation needs barriers where threads are waiting repeatedly. The parallel implementation also has overheads of context switching and maintaining cache coherence. This causes increased time in parallel implementation compared to the sequential implementation.
- 3) As we vary the number of loops, from the logarithmic graphs we can observe that the sequential and parallel implementation almost start taking the same time for loop values greater than  $10^5$  -  $10^6$ .
- 4) Loops value indicates a certain amount of work done during each operation. If the loop value is low, the sequential implementation is exponentially faster. This is because of no overhead of thread creation and barrier synchronization. As we increase the loops value, it indicates more work being done in each operation.

Since the sequential version must do this work at each of the  $O(n)$  operations, its time increases proportionally to the increase in loops value.

In the parallel case, the algorithm is still doing the  $O(n)$  order operations (work-efficient). However, there are more processing elements (PE's) which are doing this work in parallel. This implies that with increased work the parallel implementation will start yielding higher efficiency compared to a sequential implementation. Parallel algorithms still have the overhead of cache coherence and barrier synchronization which prevent them from achieving ideal speedup. However, this efficiency can be improved by increasing the number of processing elements (PE's) and better barrier primitives.

To achieve high efficiency and scalability parallel algorithms should be able to break the problem dataset into independent parts which can fit in caches. Better spatial and temporal locality in the caches can lead to higher efficiency.

3) In this step you will build your own re-entrant barrier. Recall from lecture that we considered several implementation strategies and techniques. We recommend you base your barrier on pthread's spinlocks but encourage you to use other techniques we discussed in this course. Regardless of your technique, answer the following questions: how is/isn't your implementation different from pthread barriers? In what scenarios would each implementation perform better than the other? What are the pathological cases for each? Use your barrier implementation to implement the same work-efficient parallel prefix scan. Repeat the measurements from part 2, graph them, and explain the trends in the graph. Why do these occur? What overheads cause your implementation to underperform an "ideal" speedup ratio? How do results from part 2 and part 3 compare? Are they in line with your expectations? Suggest some workload scenarios which would make each implementation perform worse than the other.

Ans –

My implementation of the spin barrier uses one spin lock and 2 atomic variables.

Pseudocode

```
spin_lock()
Incr waiters()
Incr waitingThds()

If (waiters() == expected())
{
    decrement waiters()
    do {
        // sleep for 1 ns
    } loop until waiters() == 0;
    atomic compare and exchange waiting threads = 0
    spin_unlock()
}
else
{
    spin_unlock()
    do {
        // sleep for 1 ns
    } loop until waiting threads == expected
    decrement waiters()
}
```



This implementation is different from the `pthread_barrier_t` –

- a) Each thread is constantly polling to check if the waitingThreads is equal to expected instead of waiting (scheduling out) until notified.

- b) Having a while loop waiting for the condition simulates a spin lock.

The busy wait condition (without the `nanosleep`) performs better than the `pthread_barrier` when thread count is less than the PE's. This is because it eliminates the overhead of context switching.

The sleep of 1ns was added to improve the performance compared to a busy wait for the case of higher thread count compared to processing elements (PE's).

I was observing an exponential decrease in performance without the ns sleep when the number of threads exceeds the number of process elements (PE's). This was because all the PE's were occupied with threads spinning and doing the check. Any extra threads had to wait until the current running thread is scheduled out after utilization of its time slice. Addition of the 1ns sleep helps when the number of threads exceeds the number of PE's, as this leads to scheduling out this thread and giving an opportunity for another thread to run. This provides better time performance on the barrier despite having more threads than PE's.

- c) Usage of volatile, atomic variables and CAS commands ensures that memory barriers are applied implicitly without any re-ordering.

`Pthread_barrier`'s will perform better when there are large number of threads as the barrier schedules' out each thread and puts it in a wait state. This ensures optimal utilization of the compute resources. The pathological case for `pthread_barrier`'s is having just 1 thread. This leads to lot of un-necessary work related to context switches.

My implementation using loops (without the sleep) will perform better for lesser number of threads as it will avoid the context switches overhead involved with waiting approach. The pathological case for my implementation is just having many threads (more than the number of PE's). This is because with many threads each thread must get a chance to be scheduled, to check for the validity of the condition in the while loop.

My implementation is unable to achieve the ideal speed up ratio, due to constantly polling (busy waiting) on the condition as the number of threads increase.





Please find below, the graphs - various inputs with different loop numbers from 1 – 10E7 vs execution time is plotted on a logarithmic scale of time, with different fixed number of threads. This data is from the implementation using the 1ns sleep.

## Execution time (usec) of different inputs with varying loops and fixed number of threads - custom barrier

| seq_64_test.txt |             |             |             |             |              |              | 1k.txt   |             |             |             |             |              |              |
|-----------------|-------------|-------------|-------------|-------------|--------------|--------------|----------|-------------|-------------|-------------|-------------|--------------|--------------|
|                 | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |          | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |
| 1               | 0           | 5204        | 5678        | 7514        | 11282        | 17666        | 1        | 2           | 67543       | 75872       | 101955      | 151947       | 231222       |
| 10              | 0           | 5111        | 5610        | 7525        | 11383        | 18310        | 10       | 7           | 67862       | 75684       | 102119      | 152494       | 238919       |
| 100             | 6           | 5009        | 5657        | 7535        | 11386        | 17351        | 100      | 98          | 68818       | 75681       | 100913      | 152678       | 231828       |
| 1000            | 19          | 5295        | 5657        | 7500        | 11379        | 17774        | 1000     | 317         | 71268       | 76243       | 101296      | 151237       | 230651       |
| 10000           | 185         | 5260        | 5716        | 7528        | 11458        | 17266        | 10000    | 3004        | 64355       | 76539       | 102166      | 152529       | 231025       |
| 100000          | 1837        | 7153        | 6665        | 7549        | 10811        | 16660        | 100000   | 29863       | 103895      | 96848       | 100785      | 147736       | 241421       |
| 1000000         | 18386       | 32090       | 27714       | 25302       | 26776        | 30933        | 1000000  | 298430      | 507190      | 434486      | 395261      | 376384       | 384329       |
| 10000000        | 183753      | 281463      | 237706      | 217398      | 210724       | 224571       | 10000000 | 2983840     | 4539134     | 3791887     | 3423677     | 3257409      | 3226492      |

| 8k.txt   |             |             |             |             |              |              | 16k.txt  |             |             |             |             |              |              |
|----------|-------------|-------------|-------------|-------------|--------------|--------------|----------|-------------|-------------|-------------|-------------|--------------|--------------|
| Loops    | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |          | Threads - 0 | Threads - 2 | Threads - 4 | Threads - 8 | Threads - 16 | Threads - 32 |
| 1        | 19          | 530338      | 595179      | 795172      | 1199599      | 1817060      | 1        | 39          | 1076162     | 1200775     | 1586758     | 2393754      | 3644613      |
| 10       | 55          | 538973      | 596763      | 796952      | 1198131      | 1851915      | 10       | 111         | 1069594     | 1192479     | 1595985     | 2385366      | 3574238      |
| 100      | 863         | 537213      | 608401      | 796520      | 1195260      | 1800884      | 100      | 1699        | 1078322     | 1192650     | 1584767     | 2390199      | 3629868      |
| 1000     | 2539        | 562390      | 596912      | 792735      | 1189516      | 1841030      | 1000     | 5051        | 1134219     | 1194626     | 1574738     | 2384035      | 3641775      |
| 10000    | 24079       | 509717      | 601212      | 797461      | 1184087      | 1837847      | 10000    | 48148       | 992072      | 1200598     | 1604468     | 2383671      | 4100936      |
| 100000   | 239085      | 823499      | 769921      | 785795      | 1159571      | 2000040      | 100000   | 478160      | 1646741     | 1539809     | 1567827     | 2310668      | 3941819      |
| 1000000  | 2390882     | 4050687     | 3453623     | 3158836     | 3011857      | 3048802      | 1000000  | 4778723     | 8099783     | 6908327     | 6313233     | 6062059      | 6117996      |
| 10000000 | 23890584    | 36309775    | 30351306    | 27415624    | 26905760     | 26036879     | 10000000 | 47782810    | 72609133    | 60659209    | 55492481    | 54051200     | 51632190     |



## Observations –

- 1) Just like previously seen, it can be observed that the sequential implementation seems to be much faster than parallel implementation for low values of loops. This is because the parallel implementation needs barriers where threads are waiting repeatedly. The parallel implementation also has overheads of context switching and maintaining cache coherence. This causes increased time in parallel implementation compared to the sequential implementation.
- 2) With the increase in the number of loops and threads, my barrier implementation is polling in the loops and trying to check for the condition. This leads to more work done than a `pthread_barrier` for large number of threads.

This is evident in the convergence (sequential vs parallel) loop count. We can see the sequential time converging with parallel implementation for loop counts of approximately  $10E6$  or more compared to  $10E5$  or more in the `pthread_barrier` case. (Part 2)

This is in line with the expectations of my barrier implementation.