

**Srinivas Prasad Prabhu**

**UT EID - sp55629**

**CS380P – Lab 2 – Kmeans with CUDA**

### **Platform Details**

CPU - AMD EPYC 7232P 8-Core Processor, 1.5 GHz

RAM - 128 GB RAM

OS - Ubuntu 22.04.3 LTS

GPU – Nvidia Hopper PCIe H100

41:00.0 3D controller: NVIDIA Corporation Device 233f (rev a1)

Clock – Base – 1.125, Boost - 1.755 GHz

Memory Clock – 1593 MHz

Total GPU Memory – 80GB

Memory Bus Width – 5120 Bits

Max Registers Per Block – 64k

Max Registers Per Thread – 255

Max Thread Blocks Per SM – 32

MaxThreadsPerBlock – 1024

WarpSize – 32

Concurrent warps per SM – 64

SM's per GPU – 114

Cuda Cores per SM – 128

Cuda Compute Capability – 9.0

Driver – Nvidia GPU Driver version r535.00

CUDA Toolkit – Cuda Toolkit version 12.2

NVCC Version

nvcc: NVIDIA (R) Cuda compiler driver

Copyright (c) 2005-2023 NVIDIA Corporation

Built on Tue\_Aug\_15\_22:02:13\_PDT\_2023

Cuda compilation tools, release 12.2, V12.2.140

Build cuda\_12.2.r12.2/compiler.33191640\_0

g++ (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0

gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0

Which of your implementations is fastest? Does it match your expectations of which should be fastest? Estimate the best-case performance speedup your CUDA implementations should have based on the number of threads in your program and the number of processing contexts supported by your hardware. How far of that prediction is your best-case performance?

Ans –

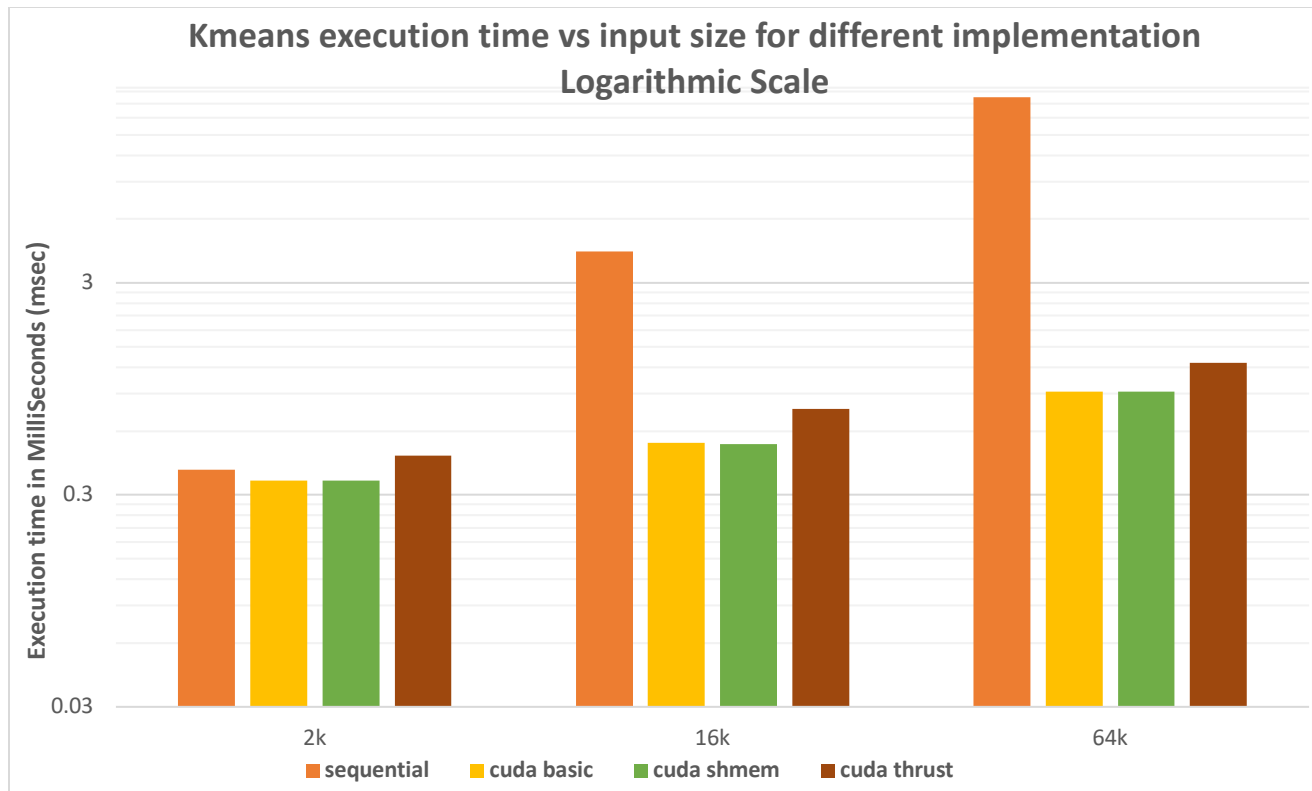
Among all my implementations the basic cuda version seems to be the fastest for inputs 16k and above. For lesser sized inputs, the difference between the cuda and sequential versions is not very significant. This can be observed from the plot below.

Note:

- 1) For cuda basic/shared versions number of threads = `maxThreadsPerBlock` = 1024
- 2) The below times were calculated using the command -

`bin/kmeans -k 16 -t 0.00000001 -d<dims> -i <file> -m 150 -s 8675309 -c -a <alg>`

Msec (per Iteration)	sequential	cuda basic	cuda shmem	cuda thrust
2k	0.393444	0.350526	0.351182	0.459445
16k	4.22568	0.526767	0.520145	0.76062
64k	22.499385	0.919002	0.91941	1.262012



Yes it matches my expectation that the CUDA version is the fastest.

The asymptotic work that needs to be done remains the same in each case.

By being able to convert the problem of kmeans from sequential to a parallel definition we can use the inherent parallelism of the GPU to complete this work in less time.

This also shows that as the input size increases, the cuda versions scale very well.

Other observations –

- 1) In sequential, cycles consumed increases exponentially with increase in input size. There is a lot of penalty paid due to constant thrashing of the caches as the algorithm keeps trying to access different points.
- 2) In my cuda implementation I have used maxThreadsPerBlock (1024) as my thread size. This utilizes a single block well; however overall occupancy is very low. This limits the amount of parallelism as the concurrent execution on an SM is limited by the warp size.

I ran some experiments so that I can increase the amount of parallelism with different blocks and thread count. I ran these experiments with 64k input and basic cuda implementation. The times below indicate that for my

implementation on my hardware, thread count of 128 gives the optimal performance.

Thread Count	Time per Iteration (msec)
1024	0.919002
512	0.959761
256	0.801223
128	0.749300
64	0.776310
32	0.770931

This experiment shows that selecting the right block, thread tuple is a carefully crafted exercise based on the implementation and hardware. **My thread selection resulted in the implementation being 18% slower on my hardware compared to the optimal output.**

Theoretically the hardware I am using has -

GPU Memory BW =  $(1593 * (10^6)) * (5120/8) * (2/(10^9)) = 2000 \text{ GB/s}$  and upto 26 TFLOPs of FP64 operations.

My kernel is only using a fraction of this available BW. On running NCU, I could see very limited occupancy and BW usage. Below output is from the profiler for 64k input.

*cudaKmeans2(options\_t \*, double \*, double \*, double \*, int \*, int \*, double)  
(64, 1, 1)x(1024, 1, 1), Context 1, Stream 7, Device 0, CC 9.0*

*Section: GPU Speed Of Light Throughput*

<i>Metric Name</i>	<i>Metric Unit</i>	<i>Metric Value</i>
--------------------	--------------------	---------------------

...		
<i>Memory Throughput</i>	<i>%</i>	<i>21.52</i>
<i>DRAM Throughput</i>	<i>%</i>	<i>0.60</i>
<i>Duration</i>	<i>msecond</i>	<i>1.37</i>

....		
<i>Compute (SM) Throughput</i>	<i>%</i>	<i>27.58</i>

*OPT This kernel grid is too small to fill the available resources on this device, resulting in only 0.5 full waves across all SMs. Look at Launch Statistics for more details.*

*Section: Launch Statistics*

<i>Metric Name</i>	<i>Metric Unit</i>	<i>Metric Value</i>
--------------------	--------------------	---------------------

<i>Block Size</i>		<i>1,024</i>
<i>Cluster Scheduling Policy</i>		<i>PolicySpread</i>
<i>Cluster Size</i>		<i>0</i>
<i>Function Cache Configuration</i>		<i>CachePreferNone</i>
<i>Grid Size</i>		<i>64</i>
<i>Registers Per Thread</i>	<i>register/thread</i>	<i>40</i>
<i>Shared Memory Configuration Size</i>	<i>Kbyte</i>	<i>8.19</i>
<i>Driver Shared Memory Per Block</i>	<i>Kbyte/block</i>	<i>1.02</i>
<i>Dynamic Shared Memory Per Block</i>	<i>byte/block</i>	<i>0</i>
<i>Static Shared Memory Per Block</i>	<i>byte/block</i>	<i>0</i>
<i>Threads</i>	<i>thread</i>	<i>65,536</i>
<i>Waves Per SM</i>		<i>0.52</i>

*OPT Estimated Speedup: 48.39%*

*The grid for this launch is configured to execute only 64 blocks,  
which is less than the GPU's 124*

### *Section: Occupancy*

<i>Metric Name</i>	<i>Metric Unit</i>	<i>Metric Value</i>
....		
<i>Max Cluster Size</i>	<i>block</i>	<i>8</i>
<i>Block Limit SM</i>	<i>block</i>	<i>32</i>
<i>Block Limit Warps</i>	<i>block</i>	<i>2</i>
<i>Theoretical Active Warps per SM</i>	<i>warp</i>	<i>32</i>
<i>Theoretical Occupancy</i>	<i>%</i>	<i>50</i>
<i>Achieved Occupancy</i>	<i>%</i>	<i>50.00</i>
<i>Achieved Active Warps Per SM</i>	<i>warp</i>	<i>32.00</i>

*cudaNewCentroids2(options\_t \*, double \*, double \*, int \*, int \*) (1, 1, 1)x(16, 1, 1), Context 1, Stream 7, Device 0, CC 9.0*

### *Section: GPU Speed Of Light Throughput*

<i>Metric Name</i>	<i>Metric Unit</i>	<i>Metric Value</i>
....		
<i>Memory Throughput</i>	<i>%</i>	<i>1.67</i>
...		
<i>SM Active Cycles</i>	<i>cycle</i>	<i>239.61</i>
<i>Compute (SM) Throughput</i>	<i>%</i>	<i>0.01</i>

*OPT This kernel grid is too small to fill the available resources on this device, resulting in only 0.0 full waves across all SMs. Look at Launch Statistics for more details.*

### *Section: Launch Statistics*

<i>Metric Name</i>	<i>Metric Unit</i>	<i>Metric Value</i>
....		
<i>Block Size</i>		<i>16</i>
<i>Grid Size</i>		<i>1</i>
<i>Registers Per Thread</i>	<i>register/thread</i>	<i>36</i>
....		
<i>Threads</i>	<i>thread</i>	<i>16</i>
<i>Waves Per SM</i>		<i>0.00</i>

- 3) Cuda shared memory implementation is not able to be better than cuda basic version. Based on my understanding, this is because
- a) We spend time initially waiting in many threads until the data is loaded into shared memory.
  - b) We are only loading the centroids into shared memory. To access the inputs, we are still going to global memory.

Only if we can reduce the number of global memory access on an aggregate considerably between the basic cuda version and the shared memory version, we will be able to see significant improvement in the shared memory time measurements.

Which of the parallel implementations is slowest, and does it match your expectations? Why or why not?

Ans -

Cuda thrust version is the slowest among all the parallel implementations. It does match my expectations as thrust is a high-level library which abstracts all the memory handling etc. in GPU's.

Typically, GPU performance is optimized by -

- a) Having good parallel algorithm
- b) Having efficient data access patterns for optimal results.

Thrust abstracts the data access patterns from the user completely. Hence can't expect thrust to provide the most optimal performance always.

Thrust provides a generic C++ STL like interface for ease of GPU programming. However, I believe under the hood it is doing additional temporary memory allocations, copies etc which adds to the time. Also, its data access patterns, placement of data into shared memory might not be optimal for a particular implementation. These factors contribute to additional time taken in the thrust implementation.

Even though thrust is not providing the best time, its ease of use and standard library of efficient algorithms makes it a good choice for quick development.



What fraction of the end-to-end runtime in your CUDA versions is spent in data transfer?

Ans-

As can be seen below, different cuda versions use less than 5% of the end-to-end time in data transfers. Most of the time is spent in actual running of the algorithm.

#### **Basic Cuda version**

Input Size	End to End Time(msec)	Algorithm Time(msec)	Difference	%age
2k	6.538848	6.309088	0.22976	3.5%
16k	13.767616	13.175296	0.59232	4.3%
64k	109.654205	107.799522	1.854683	1.7%

#### **Cuda Shared Mem version**

Input Size	End to End Time(msec)	Algorithm Time(msec)	Difference	%age
2k	6.573792	6.334016	0.239776	3.6%
16k	13.588384	12.991840	0.596544	4.3%
64k	109.619011	107.773788	1.845223	1.7%

How much time did you spend on the lab?

I spent approximately about 40 hours on this lab. It was a lot of learning. I was startled when I first saw the cuda execution time measurements. It was a learning that parallelizing algorithms can result in exponential increase in performance using GPU.